

CNCF Project Focus

Episode #4



PROMETHEUS

Cloud-Native Metrics & Monitoring



Christian Dussol



Dashboards show you what happened. **Metrics** tell you what to do next.

Most clusters have the first. Few have the second.



You can't **optimize** what you can't **measure**.



Your Kubernetes cluster runs 300 microservices.

- Which ones are over-provisioned?
- Which ones are about to break?
- How much are you actually spending?

Without **metrics**, every **decision** is a **guess**.



What you cannot see

Without metrics visibility



- ✗ No idea which services cost the most
- ✗ Over-provisioned "just in case"
- ✗ Manual capacity planning
- ✗ Budget discussions = guesswork
- ✗ No team accountability

With metrics-based visibility



- ✓ Cost attribution per namespace/pod
- ✓ Automated right-sizing alerts
- ✓ Team-level showback dashboards
- ✓ Data-driven capacity planning
- ✓ Every optimization decision backed by data

What is Prometheus ?

CNCF Graduated project (August 2018)



One of the first CNCF graduated projects
(Production-proven at hyperscale)

CNCF's first mature Observability stack foundation.

Core Idea

Pull-based metrics collection with
powerful query language

⌚ Time-series database

📊 PromQL for querying

🔔 Built-in alerting

Adopted by

NETFLIX


Uber

reddit



GitLab

From Metrics to Decisions



Prometheus for Cost Optimization

4 pillars to control your cloud spend



Resource Cost Attribution

CPU consumption per namespace

```
sum by (namespace)  
(rate(container_cpu_usage_seconds_total[5m])  
)
```



Utilization Tracking

Memory efficiency per pod

```
avg(container_memory_usage_bytes  
/ container_spec_memory_limit_bytes)
```



Waste Detection

Identify over-provisioned workloads

→ Compare actual usage vs requested resources



Budget Alerts

Trigger alerts when thresholds are exceeded

→ Combine PromQL with Alertmanager rules

Example PromQL patterns: full queries in my GitHub repo


Know your tool

Prometheus is built for

- ✓ Kubernetes metrics (native integration)
- ✓ Real-time monitoring & alerting
- ✓ Cost tracking & optimization
- ✓ SLA/SLO monitoring
- ✓ Compliance reporting (audit trails)
- ✓ Multi-tenant cost attribution

Prometheus is NOT designed for

- ✗ Full distributed tracing → Jaeger / Tempo
- ✗ Long-term scalable storage → Thanos / Cortex
- ✗ Log aggregation → Loki
- ✗ Event-based monitoring → CloudEvents



From the field

Treasury solution (Front, Middle, Back office)

Our observability stack

- 🔥 **Prometheus**: metrics & alerting
- 📊 **Grafana**: team dashboards
- 📡 **OpenTelemetry**: distributed tracing
- 🌐 **Istio**: service mesh observability

The journey

Phase 1: Deploy the stack
→ *First time we could see our system*

Phase 2: Reactive monitoring
→ *Alerts fire, investigate, fix*

Phase 3: Proactive

→ *Working with SREs to define alerts that anticipate problems before they impact users*

Deploying tools is step one.

Shifting from reactive to proactive is where it gets hard.



PROMETHEUS + KYVERNO

Observability meets Policy-as-Code



Resource Governance

Kyverno enforces CPU/memory limits on every pod
→ Prometheus metrics become accurate and meaningful



Cost Governance

Kyverno mandates team/project labels
→ Prometheus enables cost attribution per team

The Governance loop

Kyverno enforces



Prometheus measures



Teams optimize



Grafana visualizes



Educational GitHub Repository

Explore my learning toolkit



github.com/christian-dussol-cloud-native/prometheus

Complete hands-on tutorial, including:

Quick start

Cluster setup, Prometheus & Grafana installation

Complete setup scripts (Helm)

PromQL cost queries

Grafana dashboards (importable JSON)

Kyverno policies for cost governance



Educational GitHub Repository

The screenshot displays a GitHub repository interface for an educational project focused on Prometheus monitoring. The repository is titled "prometheus / prometheus-first-sample".

Repository Structure:

- Files:** README.md, LICENSE.md, README.E.md.
- Content:** Prometheus Learning - First Sample, Overview, Key Learning Objectives, Repository Organization, Quick Start Steps.

Prometheus Dashboards:

- CNCF Project Focus #4 - Cost Overview:** Shows CPU Usage by Namespace, Memory Usage by Namespace, Top 10 CPU Consumers (Pods), Top 10 Memory Consumers (Pods), CPU Usage by Team, and Memory Usage by Team.
- CNCF Project Focus #4 - Resource Efficiency:** Shows Namespaces with Over-Provisioning Indicators, CPU Utilization vs Limits (%), Memory Utilization vs Limits (%), Over-Provisioned Pods (CPU < 15% of limit), Over-Provisioned Pods (Memory < 20% of limit), Wasted CPU per Namespace (Requested - Used), and Right-Sizing Ratio (Requested / Used).

Prometheus Query Editor:

- Query:** A complex PromQL query calculating the percentage of wasted CPU across namespaces and pods.
- Table:** Results of the query, showing values for monitoring, kube-system, kyverno, and demo namespaces.
- Graph:** A graph showing the same data over time.